

Week 8 Assignment


David Russo

3/9/2017

10.7 #3

- a)

```
obs <- data.frame(  
  X1 = c(1, 1, 0, 5, 6, 4),  
  X2 = c(4, 3, 4, 1, 2, 0)  
)  
  
obs %>%  
  ggplot(aes(x = X1, y = X2)) +  
  geom_point()
```



- b)


```
# set random number seed for replicable results  
set.seed(2017)  
  
# assign the labels  
obs$cluster <- sample(rep(c(1, 2), each = 3), 6, replace = FALSE)
```

```
# print the labels
obs
```

```
##   X1 X2 cluster
## 1  1  4       2
## 2  1  3       1
## 3  0  4       1
## 4  5  1       1
## 5  6  2       2
## 6  4  0       2
```

```
# plot the labels
```

```
obs %>%
  dplyr::select(X1, X2, cluster) %>%
  ggplot(aes(x = X1, y = X2, color = as.factor(cluster))) +
  geom_point()
```



• c)

```
centroids <-
  obs %>%
  dplyr::group_by(cluster) %>%
  dplyr::summarise(centroid_X1 = round(mean(X1), 2), centroid_X2 = round(mean(X2), 2)) %>%
  as.data.frame()
```

```
centroids
```

```
##   cluster centroid_X1 centroid_X2
## 1       1         2.00         2.67
```

```
## 2      2      3.67      2.00
```

- d)

```
# create distance from cluster1 centroid
obs$dist_from_cluster1_centroid <- round(
  sqrt((obs$X1 - centroids$centroid_X1[centroids$cluster == 1])^2 +
    (obs$X2 - centroids$centroid_X2[centroids$cluster == 1])^2), 4)

# create distance from cluster2 centroid
obs$dist_from_cluster2_centroid <- round(
  sqrt((obs$X1 - centroids$centroid_X1[centroids$cluster == 2])^2 +
    (obs$X2 - centroids$centroid_X2[centroids$cluster == 2])^2), 4)

obs$new_cluster <-
  ifelse(obs$dist_from_cluster1_centroid <= obs$dist_from_cluster2_centroid, 1, 2)

obs
```

```
##   X1 X2 cluster dist_from_cluster1_centroid dist_from_cluster2_centroid
## 1  1  4       2             1.6640             3.3360
## 2  1  3       1             1.0530             2.8511
## 3  0  4       1             2.4019             4.1796
## 4  5  1       1             3.4335             1.6640
## 5  6  2       2             4.0557             2.3300
## 6  4  0       2             3.3360             2.0270
##   new_cluster
## 1             1
## 2             1
## 3             1
## 4             2
## 5             2
## 6             2
```

- e)

```
cluster_difference <- FALSE

while(cluster_difference == FALSE){

# reset obs$cluster to obs$new_cluster
obs$cluster <- obs$new_cluster

# create distance from cluster1 centroid
obs$dist_from_cluster1_centroid <- round(
  sqrt((obs$X1 - centroids$centroid_X1[centroids$cluster == 1])^2 +
    (obs$X2 - centroids$centroid_X2[centroids$cluster == 1])^2), 4)

# create distance from cluster2 centroid
obs$dist_from_cluster2_centroid <- round(
  sqrt((obs$X1 - centroids$centroid_X1[centroids$cluster == 2])^2 +
    (obs$X2 - centroids$centroid_X2[centroids$cluster == 2])^2), 4)

obs$new_cluster <-
  ifelse(obs$dist_from_cluster1_centroid <= obs$dist_from_cluster2_centroid, 1, 2)
```

```
cluster_difference <- all(obs$cluster == obs$new_cluster)

print(obs)


}
```

```
##   X1 X2 cluster dist_from_cluster1_centroid dist_from_cluster2_centroid
## 1  1  4       1             1.6640             3.3360
## 2  1  3       1             1.0530             2.8511
## 3  0  4       1             2.4019             4.1796
## 4  5  1       2             3.4335             1.6640
## 5  6  2       2             4.0557             2.3300
## 6  4  0       2             3.3360             2.0270
##   new_cluster
## 1           1
## 2           1
## 3           1
## 4           2
## 5           2
## 6           2
```

The clustering only took one iteration until the algorithm converged.

- f)

```
obs %>%
  dplyr::select(X1, X2, new_cluster) %>%
  ggplot(aes(x = X1, y = X2, color = as.factor(new_cluster))) +
  geom_point()
```



10.7 #7

```
# load the data
data("USArrests")

# scale the data so that each variable has a mean of 0 and a variance of 1
usa_scaled <- as.data.frame(scale(USArrests))
# verify that the scaled data's variables all have mean 0 and variance 1
apply(usa_scaled, 2, function(x) round(c(mean(x), var(x)), 2))
```

```
##      Murder Assault UrbanPop Rape
## [1,]      0      0      0      0
## [2,]      1      1      1      1
```

```
# create data frame of combinations of rows of usa_scaled data set
# there are choose(50, 2) = 1225 ways to choose 2 observations out of 50
combs <- t(combn(nrow(usa_scaled), 2))
head(combs)
```

```
##      [,1] [,2]
## [1,]    1    2
## [2,]    1    3
## [3,]    1    4
## [4,]    1    5
## [5,]    1    6
## [6,]    1    7
```

```
###----- Determine Correlations -----
```

```

# cycle through combs to get each correlation between rows
one_minus_cors <- rep(0, nrow(combs))

# create function for getting correlations between each observation
one_minus_cor_usa <- function(x){
  res <- cor(as.numeric(usa_scaled[combs[x, 1], ]),
             as.numeric(usa_scaled[combs[x, 2], ]))
  res <- 1 - res
  res
}

# determine 1 - r_ij for each pair of observations
one_minus_cors <- sapply(1:nrow(combs), function(x) one_minus_cor_usa(x))

###----- Determine Squared Euclidean Distances -----

# cycle through combs to get distance between rows
distances <- dist(usa_scaled)^2


###----- Determine Proportionality -----
summary(one_minus_cors/distances)

##      Min.   1st Qu.   Median     Mean   3rd Qu.     Max.
## 0.000086 0.069140 0.133900 0.234200 0.262600 4.888000

hist(one_minus_cors/distances)

```

Histogram of one_minus_cors/distances



As can be seen from the summary above, there is no evidence of proportionality. If there were proportionality, there would only be one value. I believe the wording of this question is either wrong or simply hard to understand. The question refers to distances and correlations between the i^{th} and j^{th} observations, which

would lead me to believe it is referring to the rows of the data set (i.e., 50 rows; one per state). If the authors meant to refer to the distances and correlations between i^{th} and j^{th} variables (i.e., the 4 variables of Murder, Assault, UrbanPop, and Rape), then there is constant proportionality. This is shown below.

```
# squared distance matrix
distance_matrix <- dist(t(usa_scaled))^2
distance_matrix

##           Murder  Assault UrbanPop
## Assault  19.41642
## UrbanPop 91.18188 72.63057
## Rape    42.76927 32.80636 57.68856

# 1 - correlation matrix
one_minus_cor_mat <- as.dist(1-cor(usa_scaled))
one_minus_cor_mat

##           Murder  Assault UrbanPop
## Assault  0.1981267
## UrbanPop 0.9304274 0.7411283
## Rape    0.4364212 0.3347588 0.5886588

# demonstration of proportionality
distance_matrix/one_minus_cor_mat

##           Murder Assault UrbanPop
## Assault          98
## UrbanPop         98          98
## Rape             98          98          98
```